

# edX HarvardX: PH125.8x | Capstone Project

*Nicole Yong*

*9/6/2019*

## INTRODUCTION

This Credit Card Fraud Detection project is created to fulfil the coursework requirements of the capstone module in edX HarvardX: PH125.9x (Data Science Professional Certificate).

### Credit Card Fraud Detection Dataset

This dataset contains transactions made by credit cards over two days in September 2013 by European card-holders. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

Variables:

- V1, V2, ... V28: principal components obtained with PCA
- 'Time': the seconds elapsed between each transaction and the first transaction in the dataset
- 'Amount': the transaction amount
- 'Class': the response variable; value 1 in the case of fraud and 0 otherwise

The dataset can be found on Kaggle in the following link:

- <https://www.kaggle.com/mlg-ulb/creditcardfraud/>

### Goal of Credit Card Fraud Detection Project

The goal of the credit card fraud detection project is to identify fraudulent credit card transactions.

### Measurement used: Area Under the Precision-Recall Curve

As the dataset is highly imbalanced, the accuracy will be measured using the Area Under the Precision-Recall Curve (AUPRC). The AUPRC is a single number summary of the information in the precision-recall (PR) curve.

### Key Steps Performed

#### 1. Importing data

- Downloading data from the link
- Importing data into RStudio

#### 2. Data visualization

- Explore variables in the dataset
- Draw insights from variables

### 3. Data preprocessing

- Convert target variable into factor
- Rescale variables as variables with higher values may dominate algorithms

### 4. Splitting data into train and test dataset

- Baseline train and test dataset without additional processing
- Oversampled train set: Majority Weighted Minority Oversampling Technique (MWMOTE) adjusted train dataset
- Undersampled train set: Synthetic Minority Over-sampling Technique (SMOTE) adjusted train dataset

### 5. Model selection and training

- Boosted classification trees

# METHODS & ANALYSIS OF CREDIT CARD FRAUD DETECTION DATASET

## Load Libraries

- Load the Required Libraries
- Install Missing Packages Automatically

```
## Loading required package: tidyverse
```

```
## -- Attaching packages ----- tidyverse
```

```
## v ggplot2 3.2.0      v purrr  0.3.2
## v tibble  2.1.3      v dplyr  0.8.3
## v tidyr   0.8.3      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0
```

```
## -- Conflicts ----- tidyverse_conflicts
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
## Loading required package: caret
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
## Loading required package: data.table
```

```
##
```

```
## Attaching package: 'data.table'
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
## between, first, last
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## transpose
```

```
## Loading required package: lubridate
```

```
##
```

```
## Attaching package: 'lubridate'
```

```

## The following objects are masked from 'package:data.table':
##
##     hour, isoweek, mday, minute, month, quarter, second, wday,
##     week, yday, year

## The following object is masked from 'package:base':
##
##     date

## Loading required package: ggthemes

## Loading required package: PRROC

## Loading required package: xgboost

##
## Attaching package: 'xgboost'

## The following object is masked from 'package:dplyr':
##
##     slice

## Loading required package: Matrix

##
## Attaching package: 'Matrix'

## The following object is masked from 'package:tidyr':
##
##     expand

## Loading required package: gridExtra

##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##     combine

## Loading required package: class

## Loading required package: imbalance

## Loading required package: smotefamily

## Loading required package: ROSE

## Loaded ROSE 0.0-3

##
## Attaching package: 'ROSE'

## The following object is masked from 'package:PRROC':
##
##     roc.curve

```

## Print Session Information

To help with code reproducibility, print version information about R, the OS and attached or loaded packages.

```
## R version 3.6.1 (2019-07-05)
## Platform: x86_64-mingw32/x64 (64-bit)
## Running under: Windows 10 x64 (build 18362)
##
## Matrix products: default
##
## locale:
## [1] LC_COLLATE=English_Singapore.1252 LC_CTYPE=English_Singapore.1252
## [3] LC_MONETARY=English_Singapore.1252 LC_NUMERIC=C
## [5] LC_TIME=English_Singapore.1252
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods    base
##
## other attached packages:
## [1] ROSE_0.0-3          smotefamily_1.3.1 imbalance_1.0.0
## [4] class_7.3-15        gridExtra_2.3      Matrix_1.2-17
## [7] xgboost_0.90.0.2    PRROC_1.3.1        ggthemes_4.2.0
## [10] lubridate_1.7.4     data.table_1.12.2  caret_6.0-84
## [13] lattice_0.20-38     forcats_0.4.0      stringr_1.4.0
## [16] dplyr_0.8.3         purrr_0.3.2        readr_1.3.1
## [19] tidyr_0.8.3         tibble_2.1.3       ggplot2_3.2.0
## [22] tidyverse_1.2.1
##
## loaded via a namespace (and not attached):
## [1] Rcpp_1.0.1          assertthat_0.2.1    zeallot_0.1.0
## [4] digest_0.6.20       ipred_0.9-9         foreach_1.4.4
## [7] R6_2.4.0            cellranger_1.1.0    plyr_1.8.4
## [10] backports_1.1.4     stats4_3.6.1        evaluate_0.14
## [13] httr_1.4.0          pillar_1.4.2        rlang_0.4.0
## [16] lazyeval_0.2.2      readxl_1.3.1        rstudioapi_0.10
## [19] rpart_4.1-15        rmarkdown_1.14      splines_3.6.1
## [22] gower_0.2.1         munsell_0.5.0       broom_0.5.2
## [25] compiler_3.6.1      modelr_0.1.4        xfun_0.9
## [28] pkgconfig_2.0.2     htmltools_0.3.6     nnet_7.3-12
## [31] tidyselect_0.2.5    prodlim_2018.04.18  codetools_0.2-16
## [34] crayon_1.3.4        withr_2.1.2         ModelMetrics_1.2.2
## [37] MASS_7.3-51.4       recipes_0.1.6       grid_3.6.1
## [40] nlme_3.1-140        jsonlite_1.6         gtable_0.3.0
## [43] magrittr_1.5        scales_1.0.0        cli_1.1.0
## [46] stringi_1.4.3       reshape2_1.4.3      timeDate_3043.102
## [49] xml2_1.2.0          vctrs_0.2.0         generics_0.0.2
## [52] lava_1.6.5          iterators_1.0.10     tools_3.6.1
## [55] glue_1.3.1          hms_0.5.0           survival_2.44-1.1
## [58] yaml_2.2.0          colorspace_1.4-1     rvest_0.3.4
## [61] knitr_1.23          haven_2.1.1
```

## STEP 1: DATA GATHERING & LOADING THE DATASET

### Load Credit Card Fraud Detection Dataset

```
data <- read_csv("./input/creditcard.csv")
```

Sample of dataset

```
## # A tibble: 6 x 31
##   Time      V1      V2      V3      V4      V5      V6      V7      V8      V9
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     0 -1.36 -0.0728 2.54    1.38 -0.338  0.462  0.240  0.0987  0.364
## 2     0  1.19  0.266  0.166  0.448  0.0600 -0.0824 -0.0788  0.0851 -0.255
## 3     1 -1.36 -1.34  1.77   0.380 -0.503  1.80   0.791  0.248 -1.51
## 4     1 -0.966 -0.185  1.79  -0.863 -0.0103  1.25   0.238  0.377 -1.39
## 5     2 -1.16  0.878  1.55   0.403 -0.407  0.0959  0.593 -0.271  0.818
## 6     2 -0.426  0.961  1.14  -0.168  0.421 -0.0297  0.476  0.260 -0.569
## # ... with 21 more variables: V10 <dbl>, V11 <dbl>, V12 <dbl>, V13 <dbl>,
## #   V14 <dbl>, V15 <dbl>, V16 <dbl>, V17 <dbl>, V18 <dbl>, V19 <dbl>,
## #   V20 <dbl>, V21 <dbl>, V22 <dbl>, V23 <dbl>, V24 <dbl>, V25 <dbl>,
## #   V26 <dbl>, V27 <dbl>, V28 <dbl>, Amount <dbl>, Class <dbl>
```

## STEP 2: DATA EXPLORATION & VISUALIZATION

### Part 2.1: Inspect the dataset

Summary of dataset

##	Time	V1	V2
##	Min. : 0	Min. : -56.40751	Min. : -72.71573
##	1st Qu.: 54202	1st Qu.: -0.92037	1st Qu.: -0.59855
##	Median : 84692	Median : 0.01811	Median : 0.06549
##	Mean : 94814	Mean : 0.00000	Mean : 0.00000
##	3rd Qu.: 139321	3rd Qu.: 1.31564	3rd Qu.: 0.80372
##	Max. : 172792	Max. : 2.45493	Max. : 22.05773
##	V3	V4	V5
##	Min. : -48.3256	Min. : -5.68317	Min. : -113.74331
##	1st Qu.: -0.8904	1st Qu.: -0.84864	1st Qu.: -0.69160
##	Median : 0.1799	Median : -0.01985	Median : -0.05434
##	Mean : 0.0000	Mean : 0.00000	Mean : 0.00000
##	3rd Qu.: 1.0272	3rd Qu.: 0.74334	3rd Qu.: 0.61193
##	Max. : 9.3826	Max. : 16.87534	Max. : 34.80167
##	V6	V7	V8
##	Min. : -26.1605	Min. : -43.5572	Min. : -73.21672
##	1st Qu.: -0.7683	1st Qu.: -0.5541	1st Qu.: -0.20863
##	Median : -0.2742	Median : 0.0401	Median : 0.02236
##	Mean : 0.0000	Mean : 0.0000	Mean : 0.00000
##	3rd Qu.: 0.3986	3rd Qu.: 0.5704	3rd Qu.: 0.32735
##	Max. : 73.3016	Max. : 120.5895	Max. : 20.00721
##	V9	V10	V11
##	Min. : -13.43407	Min. : -24.58826	Min. : -4.79747
##	1st Qu.: -0.64310	1st Qu.: -0.53543	1st Qu.: -0.76249
##	Median : -0.05143	Median : -0.09292	Median : -0.03276
##	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000
##	3rd Qu.: 0.59714	3rd Qu.: 0.45392	3rd Qu.: 0.73959
##	Max. : 15.59500	Max. : 23.74514	Max. : 12.01891
##	V12	V13	V14
##	Min. : -18.6837	Min. : -5.79188	Min. : -19.2143
##	1st Qu.: -0.4056	1st Qu.: -0.64854	1st Qu.: -0.4256
##	Median : 0.1400	Median : -0.01357	Median : 0.0506
##	Mean : 0.0000	Mean : 0.00000	Mean : 0.0000
##	3rd Qu.: 0.6182	3rd Qu.: 0.66251	3rd Qu.: 0.4931
##	Max. : 7.8484	Max. : 7.12688	Max. : 10.5268
##	V15	V16	V17
##	Min. : -4.49894	Min. : -14.12985	Min. : -25.16280
##	1st Qu.: -0.58288	1st Qu.: -0.46804	1st Qu.: -0.48375
##	Median : 0.04807	Median : 0.06641	Median : -0.06568
##	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000
##	3rd Qu.: 0.64882	3rd Qu.: 0.52330	3rd Qu.: 0.39968
##	Max. : 8.87774	Max. : 17.31511	Max. : 9.25353
##	V18	V19	V20
##	Min. : -9.498746	Min. : -7.213527	Min. : -54.49772
##	1st Qu.: -0.498850	1st Qu.: -0.456299	1st Qu.: -0.21172
##	Median : -0.003636	Median : 0.003735	Median : -0.06248
##	Mean : 0.000000	Mean : 0.000000	Mean : 0.00000
##	3rd Qu.: 0.500807	3rd Qu.: 0.458949	3rd Qu.: 0.13304

```

## Max. : 5.041069 Max. : 5.591971 Max. : 39.42090
## V21 V22 V23
## Min. : -34.83038 Min. : -10.933144 Min. : -44.80774
## 1st Qu.: -0.22839 1st Qu.: -0.542350 1st Qu.: -0.16185
## Median : -0.02945 Median : 0.006782 Median : -0.01119
## Mean : 0.00000 Mean : 0.000000 Mean : 0.00000
## 3rd Qu.: 0.18638 3rd Qu.: 0.528554 3rd Qu.: 0.14764
## Max. : 27.20284 Max. : 10.503090 Max. : 22.52841
## V24 V25 V26
## Min. : -2.83663 Min. : -10.29540 Min. : -2.60455
## 1st Qu.: -0.35459 1st Qu.: -0.31715 1st Qu.: -0.32698
## Median : 0.04098 Median : 0.01659 Median : -0.05214
## Mean : 0.00000 Mean : 0.00000 Mean : 0.00000
## 3rd Qu.: 0.43953 3rd Qu.: 0.35072 3rd Qu.: 0.24095
## Max. : 4.58455 Max. : 7.51959 Max. : 3.51735
## V27 V28 Amount
## Min. : -22.565679 Min. : -15.43008 Min. : 0.00
## 1st Qu.: -0.070840 1st Qu.: -0.05296 1st Qu.: 5.60
## Median : 0.001342 Median : 0.01124 Median : 22.00
## Mean : 0.000000 Mean : 0.00000 Mean : 88.35
## 3rd Qu.: 0.091045 3rd Qu.: 0.07828 3rd Qu.: 77.17
## Max. : 31.612198 Max. : 33.84781 Max. : 25691.16
## Class
## Min. : 0.000000
## 1st Qu.: 0.000000
## Median : 0.000000
## Mean : 0.001728
## 3rd Qu.: 0.000000
## Max. : 1.000000

```

Glimpse the dataset

```

## Observations: 284,807
## Variables: 31
## $ Time <dbl> 0, 0, 1, 1, 2, 2, 4, 7, 7, 9, 10, 10, 10, 11, 12, 12, 1...
## $ V1 <dbl> -1.3598071, 1.1918571, -1.3583541, -0.9662717, -1.15823...
## $ V2 <dbl> -0.07278117, 0.26615071, -1.34016307, -0.18522601, 0.87...
## $ V3 <dbl> 2.53634674, 0.16648011, 1.77320934, 1.79299334, 1.54871...
## $ V4 <dbl> 1.37815522, 0.44815408, 0.37977959, -0.86329128, 0.4030...
## $ V5 <dbl> -0.33832077, 0.06001765, -0.50319813, -0.01030888, -0.4...
## $ V6 <dbl> 0.46238778, -0.08236081, 1.80049938, 1.24720317, 0.0959...
## $ V7 <dbl> 0.239598554, -0.078802983, 0.791460956, 0.237608940, 0...
## $ V8 <dbl> 0.098697901, 0.085101655, 0.247675787, 0.377435875, -0...
## $ V9 <dbl> 0.3637870, -0.2554251, -1.5146543, -1.3870241, 0.817739...
## $ V10 <dbl> 0.09079417, -0.16697441, 0.20764287, -0.05495192, 0.753...
## $ V11 <dbl> -0.55159953, 1.61272666, 0.62450146, -0.22648726, -0.82...
## $ V12 <dbl> -0.61780086, 1.06523531, 0.06608369, 0.17822823, 0.5381...
## $ V13 <dbl> -0.99138985, 0.48909502, 0.71729273, 0.50775687, 1.3458...
## $ V14 <dbl> -0.31116935, -0.14377230, -0.16594592, -0.28792375, -1...
## $ V15 <dbl> 1.468176972, 0.635558093, 2.345864949, -0.631418118, 0...
## $ V16 <dbl> -0.47040053, 0.46391704, -2.89008319, -1.05964725, -0.4...
## $ V17 <dbl> 0.207971242, -0.114804663, 1.109969379, -0.684092786, -...
## $ V18 <dbl> 0.02579058, -0.18336127, -0.12135931, 1.96577500, -0.03...
## $ V19 <dbl> 0.40399296, -0.14578304, -2.26185710, -1.23262197, 0.80...

```



```
## $ V20 <dbl> 0.25141210, -0.06908314, 0.52497973, -0.20803778, 0.408...
## $ V21 <dbl> -0.018306778, -0.225775248, 0.247998153, -0.108300452, ...
## $ V22 <dbl> 0.277837576, -0.638671953, 0.771679402, 0.005273597, 0....
## $ V23 <dbl> -0.110473910, 0.101288021, 0.909412262, -0.190320519, -...
## $ V24 <dbl> 0.06692807, -0.33984648, -0.68928096, -1.17557533, 0.14...
## $ V25 <dbl> 0.12853936, 0.16717040, -0.32764183, 0.64737603, -0.206...
## $ V26 <dbl> -0.18911484, 0.12589453, -0.13909657, -0.22192884, 0.50...
## $ V27 <dbl> 0.133558377, -0.008983099, -0.055352794, 0.062722849, 0...
## $ V28 <dbl> -0.021053053, 0.014724169, -0.059751841, 0.061457629, 0...
## $ Amount <dbl> 149.62, 2.69, 378.66, 123.50, 69.99, 3.67, 4.99, 40.80,...
## $ Class <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0...
```

Variable names

```
## [1] "Time" "V1" "V2" "V3" "V4" "V5" "V6"
## [8] "V7" "V8" "V9" "V10" "V11" "V12" "V13"
## [15] "V14" "V15" "V16" "V17" "V18" "V19" "V20"
## [22] "V21" "V22" "V23" "V24" "V25" "V26" "V27"
## [29] "V28" "Amount" "Class"
```

Check dimension (number of rows and columns)

```
dim(data)
```

```
## [1] 284807 31
```


## Insights from Part 2.1

1. No missing data from summary table
2. All variables are masked, except 'Amount', 'Time' and 'Class'

## Part 2.2: Inspect target variable - 'Class'

```
table(data$Class)
```

```
##  
##      0      1  
## 284315  492
```



Convert target variable ('Class') into factor

```
data$Class <- factor(ifelse(data$Class == 0, "zero", "one"))
```

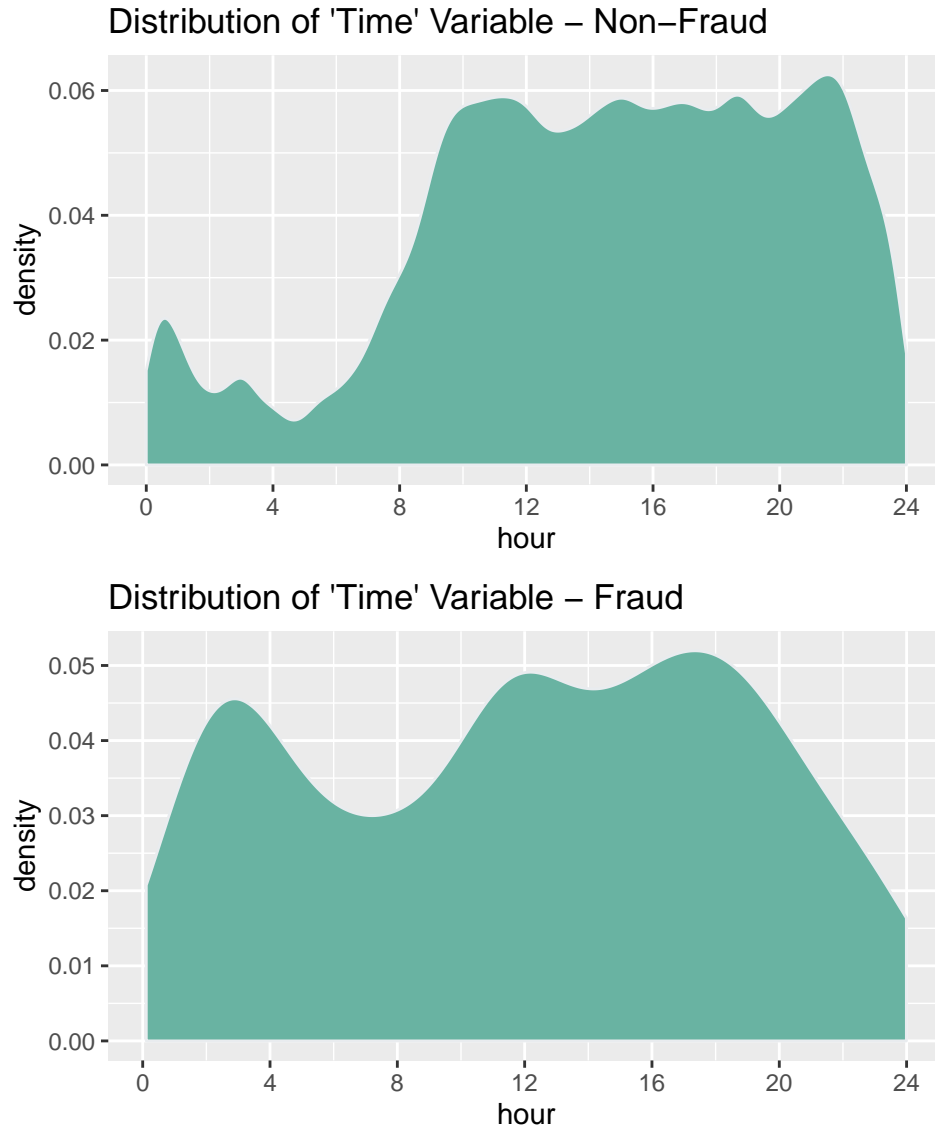
## Insights from Part 2.2

1. Highly imbalanced dataset

### Part 2.3: Inspect 'Time' variable

Convert seconds to hours of the day


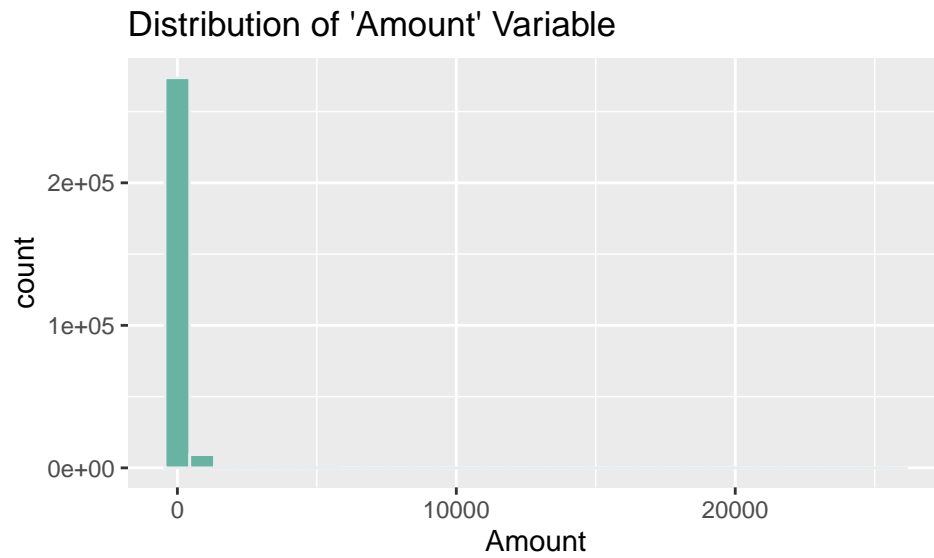
```
data$hour <- (data$Time/3600) %% 24
```





### Insights from Part 2.3

1. Non-fraud credit card transactions started to pick up from 08:00AM, maintained traction, and fell off at roughly 22:00PM
2. This would be in line with how a normal person would be using his/her credit card; starting from breakfast hours, throughout the day, and tapering off when the day ends
3. Fraudulent credit card transactions had more activity during 00:00 midnight to 08:00AM
4. This could be because credit card transactions made during these hours, when the real owners are presumably asleep, are less likely to be found out via real-time bank alerts etc.
5. This could also suggest fraudulent credit card transactions are made in a different time zone

#### Part 2.4: Inspect 'Amount' variable



- 'Amount' variable cannot be seen properly without log transformation
- Hence the log transformed 'Amount' variable will be used in the charts below



#### Insights from Part 2.4

1. The summary table shows the min. for 'Amount' variable is 0.00, which seems counter-intuitive to a credit card transaction
2. Potentially, a 0.00 amount could be used to verify if the credit card is a valid one
3. Fraudulent transactions have amounts smoothened out with less variability in the distribution
4. This may suggest certain 'favored' amounts that credit card fraudsters charge to the credit cards
5. This is unlike non-fraudulent transactions, where credit card transactions can be for a wide range of amounts since transactions for goods and services are unlimited to certain amounts

## STEP 3: DATA PREPROCESSING

### Part 3.1: Rescale variables to be in the range 0 to 1

- Variables with higher values may dominate computations and skew the model performance
- Hence, normalize data to work with different variables that are in different scales

```
##           min           max
## Time      0.000000 1.727920e+05
## V1      -56.407510 2.454930e+00
## V2      -72.715728 2.205773e+01
## V3      -48.325589 9.382558e+00
## V4       -5.683171 1.687534e+01
## V5     -113.743307 3.480167e+01
## V6      -26.160506 7.330163e+01
## V7      -43.557242 1.205895e+02
## V8      -73.216718 2.000721e+01
## V9      -13.434066 1.559499e+01
## V10     -24.588262 2.374514e+01
## V11      -4.797473 1.201891e+01
## V12     -18.683715 7.848392e+00
## V13      -5.791881 7.126883e+00
## V14     -19.214325 1.052677e+01
## V15      -4.498945 8.877742e+00
## V16     -14.129855 1.731511e+01
## V17     -25.162799 9.253526e+00
## V18      -9.498746 5.041069e+00
## V19      -7.213527 5.591971e+00
## V20     -54.497720 3.942090e+01
## V21     -34.830382 2.720284e+01
## V22     -10.933144 1.050309e+01
## V23     -44.807735 2.252841e+01
## V24      -2.836627 4.584549e+00
## V25     -10.295397 7.519589e+00
## V26      -2.604551 3.517346e+00
## V27     -22.565679 3.161220e+01
## V28     -15.430084 3.384781e+01
## Amount    0.000000 2.569116e+04
## hour       0.000000 2.399944e+01

##      min_after_rescaling max_after_rescaling
## Time                    0                    1
## V1                      0                    1
## V2                      0                    1
## V3                      0                    1
## V4                      0                    1
## V5                      0                    1
## V6                      0                    1
## V7                      0                    1
## V8                      0                    1
## V9                      0                    1
## V10                     0                    1
## V11                     0                    1
## V12                     0                    1
```

## V13	0	1
## V14	0	1
## V15	0	1
## V16	0	1
## V17	0	1
## V18	0	1
## V19	0	1
## V20	0	1
## V21	0	1
## V22	0	1
## V23	0	1
## V24	0	1
## V25	0	1
## V26	0	1
## V27	0	1
## V28	0	1
## Amount	0	1
## hour	0	1

## STEP 4: SPLIT DATASET INTO TRAIN & TEST SET

Set seed

```
set.seed(1)
```

### Part 4.1: Baseline dataset

Train data: 70%; Test data: 30%

```
index_train <- createDataPartition(data$Class, p=0.7, list=FALSE)

baseline_train <- data[index_train, ]
test <- data[-index_train, ]

baseline_train$Class <- factor(baseline_train$Class)
```



## Part 4.2: MWMOTE adjusted dataset

- MWMOTE: Majority Weighted Minority Oversampling Technique
- Oversampling helps to balance class distribution by duplicating minority class instances

Generate more noisy instances out of dataset

```
mwmote_train_gen <- mwmote(as.data.frame(baseline_train), numInstances=100000)
table(mwmote_train_gen$Class)
```

```
##
##      one    zero
## 100000      0
```

Bind the train dataset and the newly generated instances

```
mwmote_train <- rbind(baseline_train, mwmote_train_gen)
```

Compare 'Class' variable in the baseline dataset and MWMOTE adjusted dataset

```
table(baseline_train$Class)
```

```
##
##      one    zero
##    345 199021
```

```
table(mwmote_train$Class)
```

```
##
##      one    zero
## 100345 199021
```

```
mwmote_train$Class <- factor(mwmote_train$Class)
```

## Insights from Part 4.2

1. MWMOTE adjusted dataset is a more balanced dataset

### Part 4.3: SMOTE adjusted dataset

- Undersample instances of the majority class so classifiers are not biased toward one class

```
set.seed(1)
smote_train_gen1 <- SMOTE(X = baseline_train[, -1], target = baseline_train$Class, dup_size = 4)
smote_train_gen2 <- smote_train_gen1$data %>% rename(Class = class)
```

Undersample dataset until majority class size matches

```
smote_train_gen3 <- ovun.sample(Class ~ .,
                                data = smote_train_gen2,
                                method = "under",
                                N = 2 * sum(smote_train_gen2$Class == "one"))

smote_train <- smote_train_gen3$data
```

Compare 'Class' variable in the baseline dataset and SMOTE adjusted dataset

```
table(baseline_train$Class)
```

```
##
##   one   zero
##  345 199021
```

```
table(smote_train$Class)
```

```
##
##   one zero
## 1725 1725
```

```
smote_train$Class <- factor(smote_train$Class)
```

### Insights from Part 4.3

1. SMOTE adjusted dataset is a more balanced dataset

## STEP 5: MODEL SELECTION & TRAINING

- Evaluation Metric: Area Under the Precision-Recall Curve (AUPRC)

### XGBoost: xgbTree Method

- Build a series of trees where each tree is trained to attempt to correct the mistakes of the previous tree in the series
- To create a model that makes fewer and fewer mistakes as more trees are added
- Making predictions with gradient boosted tree models is faster

#### Part 5.1: XGBoost model for baseline dataset

```
##      nrounds max_depth eta gamma colsample_bytree min_child_weight subsample
## 32      100         2 0.3    0              0.8              1         0.75

## [1] "Area Under the Precision-Recall Curve: 0.873"
```

Add baseline model results to a table for later comparison

Model	AUPRC
Baseline XGBoost	0.873

## Part 5.2: XGBoost model for MWMOTE adjusted dataset

```
##      nrounds max_depth eta gamma colsample_bytree min_child_weight subsample
## 54      150         3 0.3    0             0.8             1             1
```

```
## [1] "Area Under the Precision-Recall Curve: 0.86"
```

Add baseline model results to a table for later comparison

Model	AUPRC
Baseline XGBoost	0.873
MWMOTE Adjusted XGBoost	0.860

### Part 5.3: XGBoost model for SMOTE adjusted dataset

```
##      nrounds max_depth eta gamma colsample_bytree min_child_weight subsample
## 51      150          3 0.3      0              0.8              1      0.75
```

```
## [1] "Area Under the Precision-Recall Curve: 0.722"
```

Add baseline model results to a table for later comparison

Model	AUPRC
Baseline XGBoost	0.873
MWMOTE Adjusted XGBoost	0.860
SMOTE Adjusted XGBoost	0.722

# RESULTS

## Modeling Results & Model Performance

To recap, the goal of this project is to identify fraudulent credit card transactions. As the dataset is highly imbalanced, the accuracy will be measured using the Area Under the Precision-Recall Curve (AUPRC).

The modeling results and model performance shows that the baseline training dataset delivered better model performance. In summary, the highest AUPRC was **0.873**, achieved using **XGBoost** with **xgbTree** method.

This reflects the challenges in dealing with a highly imbalanced dataset. An examination of the XGBoost model results shows that the *max\_depth* (which controls the depth of the tree) using the baseline train dataset was **2**, while using the MWMOTE adjusted and SMOTE adjusted train datasets returned a *max\_depth* of **3**. This could potentially explain why the baseline trained model scored the highest AUPRC. It is possible that both models trained with the adjusted datasets were overfitting.

# CONCLUSION

## Brief Summary

In summary, this report explained the steps in creating a machine learning algorithm to detect fraudulent credit card transactions from the Credit Card Fraud Detection dataset (from Kaggle). This report also dealt with a highly imbalanced dataset and included trying out various preprocessing steps for such datasets. The boosting method from the **XGBoost** package is used for this classification problem, with Area under the Precision-Recall curve of **0.873**.

## Limitations and Future Work

For the purpose of this project, the Credit Card Fraud Detection dataset provided by Kaggle had most of the variables masked via PCA. As such, while the algorithms may accurately detect a fraudulent credit card transaction, it does not provide a rational explanation as to why a transaction was categorized as fraudulent. For future work, with unmasked variables, we may use the machine learning algorithm to better understand which are the variables that would lead to a transaction being classified as fraudulent, and may further work towards this to reduce credit card frauds.

## Github

The following link to the GitHub repository contains the reports in PDF format, Rmd format and R script: <https://github.com/Nicole-Yong/Credit-Card-Fraud-Analysis>